Constructing Parsimonious Hybridization Networks with Multiple Phylogenetic Trees using SAT-solver: Initial Explorations

Vladimir Ulyantsev and Mikhail Melnik

ITMO University, Saint-Petersburg, Russia
{ulyantsev,melnik}@rain.ifmo.ru

Abstract. TODO: abstract

Keywords: TODO

1 Introduction

Phylogenetic network is a powerful tool to model the reticulate evolutionary processes (such as horizontal gene transfer and hybrid specification). Briefly, phylogenetic network is a directed acyclic graph, which has nodes (called reticulation nodes) with more than one incoming edge. Phylogenetic networks have been studied by several (?) researchers [5], [6], [7]. There are several formulations of phylogenetic network construction problem with various modeling assumptions and different types of input data. In this paper we focus on the specific type of phylogenetic network called hybridization network [8], [3].

Input data for hybridization network construction is a set of gene trees. Each gene tree models the evolutionary history of some gene. Because of the reticulate evolutionary events these trees can have different topologies. The aim is to construct a phylogenetic network containing the smallest possible number of reticulation nodes and displaying each of the input trees.

Most of the algorithms for phylogenetic network construction are heuristic [??] and they usually deal with only two trees. However, recently an exact algorithm PIRN has been intoduced [10] which is able to process more than two input trees.

In this paper we introduce a new approach for parsimonius phylogenetic network construction from multiply input trees based on satisfiabilty (SAT) solvers. SAT problem is known to be NP-complete [2], but state-of-the-art SAT-solvers running on the modern hardware are able to solve SAT instances having tens of thousands of variables and hundreds of thousands of clauses in several minutes.

SAT-solver based algorithms have been successfully applied to efficiently calculate evolutionary tree measures [1] as well as to solve problems in other domains such as finite-state machine induction [4], software verification [], ...

The general outline of our approach is to convert an instance of phylogenetic network construction problem to an instance of SAT problem (a Boolean formula), solve it using SAT-solver and then convert the satisfying assignment into

the phylogenetic network. The formula construction process should ensure that the network has the minimal possible number of reticulate nodes.

Our approach leads to an exact algorithm outperforming PIRN on ??? test-cases out of ???

The paper is structured as follows. Section 2 gives the formal definitions, section 3 describes the Boolean formula construction process and section 4 gives the experimental results.

2   Definitions and Background

We define a "phylogenetic tree" as a leaf-labeled tree constructed over a set of taxa. Throughout this paper we assume that trees are rooted and binary, i.e. in-degrees of all nodes, except root, are one, and out-degrees are zero for leaves and two for internal nodes. A "hybridisation network" is a directed acyclic graph with a single root and leaves bijectively labeled by the set of taxa. If the in-degree of a node is greater than one, such node is called hybridisation node. In this paper we assume that hybridisation nodes have in-degree of two and out-degree of one. We can do this assumption by noting that we can convert a hybridisation node with in-degree of three or more to the sequence of hybridisation nodes of in-degree two [9]. Other nodes are regular tree nodes and have the same properties as in tree.

Consider a hybridisation network N. Firstly we keep only one of the incoming edges for each hybridisation node in N. Secondly we contract edges to remove all the nodes of in-degree one and out-degree one. With this two steps we obtain a tree T' from the network N.

Now suppose we have a phylogenetic tree T and a hybridisation network N. We say that network N *displays* T if we can choose the edges of hybridisation nodes in such a way that after edge-contraction obtained tree T' will be equivalent to T.

A "hybridisation number" is commonly defined as the sum of in-degrees of all edges minus one. Note that under our assumptions $h(N)$ simply equals the number of reticulation nodes.

$h(N) = \sum\limits_{v!=root} (d^-(v) - 1)$

Suppose we are given a set of K phylogenetic trees $T_1, T_2, ...T_k$ over the same set of taxa. The minimum hybridisation network for that set of trees is a network N_{min} that displays each tree and has the lowest hybridisation number possible.

The most parsimonious hybridisation network problem.

Given a set of K phylogenetic trees $T_1, T_2, ...T_k$ over the same set of taxa, construct the minimum hybridisation network for that set of trees.

It has been shown that even for the case of $K = 2$ construction of such a network is NP-complete problem [2]. As far as we know there exists only one algorithm for the construction of the most parsimonious hybridisation network [10].

3 Algorithm

The main idea of the algorithm is to look over possible values if hybridisation number and to construct a boolean formula that represent a hybridisation network with fixed hybridisation number. Before doing this search we do some preprocessing and consequently some postprocessing after the network is found.

3.1 Preprocessing

Before the actual boolean formula encoding we modify the input and split it into several tasks to reduce the size and complexity of the problem. We do this according to the rules from [1]. To define those rules we first need to define the term "cluster". A set of taxa A is a cluster in trees $T_1, T_2, ...T_k$ if there exists a node in each tree that has all the taxa from A in its descendants.

Subtree reduction rule

Replace any subtree that appears identically in all input trees by a leaf with new label.

Cluster reduction rule

For each cluster A replace the subtrees containing it by a leaf with new label and add a new task for processing that consists of deleted subtrees $T'_1, T'_2, ...T'_k$ with leaf set A.

After we split the task into a set of less complex tasks, we add a dummy root to each tree of each task along with the new dummy leaf for tree consistency. This is done to ensure that all the trees in the input share common root. This dummy root will be deleted on the postprocessing stage. At this stage we have a set of tasks that will be solved separately and then their results will be merged at the postprocessing stage.

3.2 Search of the minimal hybridisation number

To solve a subtask we need to find the lowest hybridisation number k such that there exists a hybridisation network with this hybridisation number. We look through possible values of k starting from the highest (?? why do we choose max k equals to the size of taxa ??) and construct a boolean formula corresponding to that k. We decrease k until the solver can't satisfy the formula, and that means that the previous value of k was the lowest possible. There are another strategies of searching the minimal value of k. For example we can start from zero and increase k until the solver will be able to satisfy the formula, or we can use binary search. But in our experiments the most time-consuming computations were with the value of $k_{optimal} - 1$ and consequently results of upwards method were very poor. Binary search results were close to downwards method, but in some cases when binary search tried to satisfy formulas with values of k less than $k_{optimal}$ its results also were poor. This can be explained by an observation that it is much easier to solver to produce an answer if formula is satisfiable than if it is not. In case of unsatisfiable formula solver should check every possible answer to ensure that there is no solution. Obvious method for reducing searching

time is to use different heuristics for finding close starting upper and lower bounds of hybridisation number. Possible candidates are PIRN, RIATA-HGT and MURPAR. We do not consider this optimisations in this paper.

3.3   Encoding boolean formula

Having a set of trees T over a set of taxa N and fixed hybridisation number k we will construct boolean formula that will be satisfiable iff there exists a hybridisation network that displays each tree in T and its hybridisation number equals to k. To do this we first notice that a network over the set of taxa of size n with hybridisation number k will have $2 * (n + k) - 1$ nodes (?? why ??). As we add dummy root and dummy leaf, we finally have $2 * (n + k) + 1$ nodes, k of which are hybridisation nodes, $n + 1$ are leaves and others are usual tree nodes. Let R be the set of hybridisation nodes, L be the set of leaves and V be the set of regural nodes. We number all the nodes in such a way that all leaves have numbers in range $[0, n]$, all regular nodes have numbers in range $[n + 1, 2n + k]$ and all hybridisation nodes have numbers in range $[2n + k + 1, 2(n + k)]$. And we assume that the number of any leave or regular node is less than number of its parent. Such numeration allows us to define the following sets of nodes for each node v: $PC(v)$ is the set of possible children of v, $PP(v)$ is the set of possible parents of v and $PU(v)$ is the set of nodes that can be indirect parents of v. Now we will describe literals and clauses required to construct boolean formula in conjunctive normal form.

Network structure encoding First of all we encode structure of the network itself. We introduce the following literals:

- $\forall v \in V, u \in PC(v) : l_{v,u}(r_{v,u})$
  $l(r)$ is true iff regular node v has node u as left (right) child.
- $\forall v \in L \cup V/root, u \in PP(v) : p_{v,u}$
  p is true iff u is parent of regular node v.
- $\forall v \in R, u \in PP(v) : lp_{v,u}(rp_{v,u})$
  $lp(rp)$ is true iff u is left (right) parent of hybridisation node v.
- $\forall v \in R, u \in PC(v) : ch_{v,u}$
  ch is true iff u is child of hybridisation node v.

This takes $O((n + k)^2)$ literals for specifying network structure, and if we notice that $k < n$ then it is equal to $O(n^2)$ literals.

To encode the uniqueness of parents and children we use an obvious pattern that requires $O(n^2)$ clauses. We say that the node can have at least one real parent (child or something) and at most one true parent. This can be done in the following way (for example for parents of regular node v):

$$\left(\bigvee_{u \in PP(v)} p_{v,u} \right) \wedge \left(\bigwedge_{i,j \in PP(v); i < j} (p_{v,i} \to \neg p_{v,j}) \right).$$

We introduce the following constraints to restrict the network structure:

- at-most-one and at-least-one clauses for literals l, r, p, lp, rp, ch

– left child must have less number then the right one. This constraint can be translated as
$\forall u, w \in PC(v) : u \geq w$
$l_{v,u} \rightarrow \neg r_{v,w}$
– same constraint for left and right parents of hybridisation nodes.

We should connect parent literals with children literals for network consistency. We do this with the following constraints:

– if regular node has regular children then children should have this node as parent and vice versa:
  • if u is the left child of v then v should be the parent of u
    $left_{v,u} \rightarrow parent_{u,v}$
  • the same for the right child
    $right_{v,u} \rightarrow parent_{u,v}$
  • if v is the parent of u then v is either its right or left child
    $parent_{u,v} \rightarrow (left_{v,u} \vee right_{v,u})$

The same idea applies also for other types of parent-child relations.

– regular parents with hybridisation children
  • $left_{v,u} \rightarrow (lp_{u,v} \vee rp_{u,v})$
  • $right_{v,u} \rightarrow (lp_{u,v} \vee rp_{u,v})$
  • $lp_{u,v} \rightarrow (left_{v,u} \vee right_{v,u})$
  • $rp_{u,v} \rightarrow (left_{v,u} \vee right_{v,u})$
– hybridisation parents with regular children
  • $ch_{v,u} \rightarrow parent_{u,v}$
  • $parent_{u,v} \rightarrow ch_{v,u}$
– hybridisation parents with hybridisation children
  • $ch_{v,u} \rightarrow (lp_{u,v} \vee rp_{u,v})$
  • $lp_{u,v} \rightarrow ch_{v,u}$
  • $rp_{u,v} \rightarrow ch_{v,u}$
– we also should order parents and children of hybridisation nodes, for consistent numeration
  $\forall v, u, w : v \in R, u \in PC(v), w \in PP(v) : w \leq u$
  • $ch_{v,u} \rightarrow \neg lp_{v,w}$
  • $ch_{v,u} \rightarrow \neg rp_{v,w}$

Mapping trees to network To show that network should contain all the input trees we add literals that represent mapping the tree over the network. We should note that leaves of the trees are bijectively mapped to leaves of the network. (?? The intution for these literals is that ??):

– $\forall t \in T, s \in V(t), v \in V : x_{t,s,v}$
  x is true iff regular node v represents node s from tree t, i.e. x literals represent bijective mapping of network nodes to tree nodes

- $\forall t \in T, v \in R : dir_{t,v}$
  dir is true iff left parent edge of hybridisation node v is used to display tree t and it is false in case of right edge, i.e. it specifies direction of necessary parent to display current tree
- $\forall t \in T, v \in R : rused_{t,v}$
  $rused$ is true iff child of hybridisation node v is used to display tree t
- $\forall t \in T, v \in V : used_{t,v}$
  $used$ is true iff regular node v is used to display tree t
- $\forall t \in T, v \in V, u \in PU(v) : up_{t,v,u}$
  up is true iff regular node u is indirect parent of node v in the network but it is its direct parent in tree t

We use up variables because edge contraction is used for displaying trees, so it is not enough to only consider direct parents in network.

Now we define constraints that show whether network node is used to display an input tree:

- First of all we forbid meaningless values of $used$ literals.
  If the direction of parent of hybridisation node for specific tree is defined, we shouldn't use another branch:
  - $\forall t \in T, v \in R, u \in PP(v) :$
    $(\neg dir_{t,v} \wedge lp_{v,u}) \rightarrow \neg used_{t,u}$
    $(dir_{t,v} \wedge rp_{v,u}) \rightarrow \neg used_{t,u}$
  Also if we don't use current hybridisation node for current tree, we shouldn't use its parents for current tree:
  - $\forall t \in T, v \in R, u \in PP(v) :$
    $(lp_{v,u} \wedge \neg rused_{t,v}) \rightarrow \neg used_{t,u}$
    $(rp_{v,u} \wedge \neg rused_{t,v}) \rightarrow \neg used_{t,u}$
- Secondly we limit usage of hybridisation nodes.
  $\forall t \in T, v \in R, u \in PP(v) :$
  If $u \in V \cup L$, i.e. u is regular tree node. Then we have no choice but to use hybridisation node:
  - $ch_{t,v,u} \rightarrow rused_{t,v}$
  Then consider $u \in R$.
  If child is not used for current tree then we shouldn't use the node itself:
  - $\neg rused_{t,u} \rightarrow \neg rused_{t,v}$
  We also should forbid usage of current node when its direction differs from direction specified in its child, and we should use it if its direction matches direction specified in its child and its child itself is used:
  - $(lp_{t,u,v} \wedge \neg dir_{t,u}) \rightarrow \neg rused_{t,v}$
  - $(rp_{t,u,v} \wedge dir_{t,u}) \rightarrow \neg rused_{t,v}$
  - $(lp_{t,u,v} \wedge dir_{t,u} \wedge rused_{t,u}) \rightarrow rused_{t,v}$
  - $(rp_{t,u,v} \wedge \neg dir_{t,u} \wedge rused_{t,u}) \rightarrow rused_{t,v}$
- At-least-one and at-most one clauses for literals x and up. Note that in case of x literals we have restriction that at most one node from tree corresponds to the network node, and that at most one node from network corresponds to the tree node.

– If we know that node from tree bijectively matches some network node then we should use that node for that tree:

- $x_{t,s,v} \rightarrow used_{t,s}$

Note that due to dummy roots we know that all tree roots match network root:

- $x_{t,s,root}$

– Constraints for regular node v and its relation to direct tree parent from tree t.

If parent u of node v is used in tree t, then it is its direct parent in corresponding tree and vice versa:

- $\forall t \in T, v \in V \cup L, u \in PP(v), u \in V$
  $(parent_{v,u} \wedge used_{t,u}) \rightarrow up_{t,v,u}$
  $(parent_{v,u} \wedge up_{t,v,u}) \rightarrow used_{t,u}$

If we do not use parent u of node v in tree t, we should traverse direct tree parent from u to v and vice versa:

- $\forall t \in T, v \in V \cup L, u \in PP(v), u \in V, w \in PP(u)$
  $(parent_{v,u} \wedge \neg used_{t,u} \wedge up_{t,u,w}) \rightarrow up_{t,v,w}$
  $(parent_{v,u} \wedge \neg used_{t,u} \wedge up_{t,v,w}) \rightarrow up_{t,u,w}$

If parent u of node v is hybridisation node, we should care about order of v and parents of u:

- $\forall t \in T, v \in V \cup L, u \in PP(v), u \in R, w \in PU(u)$
  if $w \leq v$ then w can not be parent of u considering tree t:
  $parent_{v,u} \rightarrow \neg up_{t,u,w}$
  otherwise we should traverse direct tree parents from u to v (?? and maybe vice versa ??):
  $(parent_{v,u} \wedge up_{t,u,w}) \rightarrow \neg up_{t,v,w}$

– Constraints for hybridisation node v and its relation to direct tree parent from t.

If direction of parent u matches direction specified in node v and parent is used in t then is is direct parent of v in t:

- $\forall t \in T, v \in R, u \in PP(v), u \in V$
  $(lp_{v,u} \wedge dir_{t,v} \wedge used_{t,u}) \rightarrow up_{t,v,u}$
  $(rp_{v,u} \wedge \neg dir_{t,v} \wedge used_{t,u}) \rightarrow up_{t,v,u}$

If parent u us not used in tree t then we should traverse direct tree parents from u to v (?? and maybe vice versa ??):

- $\forall t \in T, v \in R, u \in PP(v), u \in V, w \in PU(u)$
  $(lp_{v,u} \wedge dir_{t,v} \wedge \neg used_{t,u} \wedge up_{t,u,w}) \rightarrow up_{t,v,w}$
  $(rp_{v,u} \wedge \neg dir_{t,v} \wedge \neg used_{t,u} \wedge up_{t,u,w}) \rightarrow up_{t,v,w}$

If parent u is a hybridisation node we also should traverse direct tree parents from u to v (?? and maybe vice versa ??):

- $\forall t \in T, v \in R, u \in PP(v), u \in R, w \in PU(u)$
  $(lp_{v,u} \wedge dir_{t,v} \wedge up_{t,u,w}) \rightarrow up_{t,v,w}$
  $(rp_{v,u} \wedge \neg dir_{t,v} \wedge up_{t,u,w}) \rightarrow up_{t,v,w}$

- Consider restrictions on x literals.
  If v is leaf and its parent u bijectively maps to node tu in tree t then it is direct parent.
  - $\forall t \in T, v \in L, u \in PP(v), tu \in V(t)$
    $x_{t,tu,u} \rightarrow up_{t,v,u}$

  If node v and its parent u both bijectively map to child and parent in tree t, then u is direct parent of v considering tree t.
  - $\forall t \in T, v \in V, u \in PP(v), tv \in V(t), tu = p(tv)$
    $(x_{t,tv,v} \wedge x_{t,tu,u}) \rightarrow up_{t,v,u}$

  On the other hand if we know that u is direct parent of v in tree t, then we also have bijective mapping for them.
  - $\forall t \in T, v \in V, u \in PP(v), tv \in V(t), tu = p(tv)$
    $(x_{t,tv,v} \wedge up_{t,v,u}) \rightarrow x_{t,tu,u}$

  We also should care about numeration. Thus we can not map node u from network to node tu from tree if number of u is less than number of its child v.
  - $\forall t \in T, v \in V, u \in PP(v), tv \in V(t), tu = p(tv), u < v$
    $x_{t,tv,v} \rightarrow \neg x_{t,tu,u}$

Heuristics

- We add some heuristic constraints by noticing an obvious fact that number of node in network can not be less than the size of its subtree and it can not be greater than $|t| - depth(tv)$.
  - $\forall t \in T, v \in V, tv \in V(t), tv < |subtree(tv)|$
    $\neg x_{t,tv,v}$
  - $\forall t \in T, v \in V, tv \in V(t), tv > |t - depth(tv)|$
    $\neg x_{t,tv,v}$
- Consider node tv_1 in tree t_1 and node tv_2 in tree t_2 such that thier subtrees have disjoint set of taxa. Then these nodes can not be mapped to the same node in network.
  - $\forall t_1, t_2 \in T, v \in V, tv_1 \in V(t_1), tv_2 \in V(t_2)$
    $\neg x_{t_1,tv_1,v} \vee \neg x_{t_2,tv_2,v}$

3.4   Solving boolean formula

To solve generated boolean formula we use a SAT-solver CryptoMiniSat 4.2.0. Choosing the most appropriate solver is not considered in current paper, however several experimentations were made by M. Bonet and K. John [1] so it is one of the topics of further research.

3.5   Postprocessing

After solving a task we recover network from the SAT-solver output. After that we delete dummy root and appropriate leaf from the network. And when all the subtasks of the original task are solved, we merge their networks into one hybridisation network corresponding to original task.

4   Results

To test performance of our algorithm we evaluated it on a grass (Poaceae) dataset provided by the Grass Phylogeny Working Group (Grass Phylogeny Working Group, 2001). We compared results our method with results of PIRN in exact and in heuristic modes.

5   Discussion

Testing shows that our method outperform exact PIRN algorithm in all cases.

Acknowledgements

References

1. Bonet, M.L., John, K.S.: Efficiently calculating evolutionary tree measures using sat. In: Theory and Applications of Satisfiability Testing-SAT 2009, pp. 4–17. Springer (2009)
2. Bordewich, M., Semple, C.: Computing the minimum number of hybridization events for a consistent evolutionary history. Discrete Applied Mathematics 155(8), 914–928 (2007)
3. Chen, Z.Z., Wang, L.: Hybridnet: a tool for constructing hybridization networks. Bioinformatics 26(22), 2912–2913 (2010)
4. Heule, M.J., Verwer, S.: Exact dfa identification using sat solvers. In: Grammatical Inference: Theoretical Results and Applications, pp. 66–79. Springer (2010)
5. Huson, D.H., Rupp, R., Scornavacca, C.: Phylogenetic networks: concepts, algorithms and applications. Cambridge University Press (2010)
6. Morrison, D.A.: Introduction to phylogenetic networks. RJR Productions (2011)
7. Nakhleh, L.: Evolutionary phylogenetic networks: models and issues. In: Problem solving handbook in computational biology and bioinformatics, pp. 125–158. Springer (2011)
8. Semple, C.: Hybridization networks. Department of Mathematics and Statistics, University of Canterbury (2006)
9. Wu, Y.: Close lower and upper bounds for the minimum reticulate network of multiple phylogenetic trees. Bioinformatics 26(12), i140–i148 (2010)
10. Wu, Y.: An algorithm for constructing parsimonious hybridization networks with multiple phylogenetic trees. Journal of Computational Biology 20(10), 792–804 (2013)